

Nearest Neighbour Algorithms

Christos Dimitrakakis

September 19, 2025

Outline

Introduction

The hidden secret of machine learning

The algorithm

k Nearest Neighbours

Extensions and parameters

Activities

Introduction

The hidden secret of machine learning

The algorithm

k Nearest Neighbours

Extensions and parameters

Activities

Fill-in class data



Figure: Spreadsheet link

(Batch) learning from data

- ▶ $x_t \in \mathcal{X}$, the **input** variables (or **features**)
- ▶ Usually $\mathcal{X} = \mathbb{R}^n$: the n-dimensional Euclidean space

(Batch) learning from data

- ▶ $x_t \in \mathcal{X}$, the **input** variables (or **features**)
- ▶ Usually $\mathcal{X} = \mathbb{R}^n$: the n -dimensional Euclidean space

Unsupervised learning:

- ▶ Training examples (x_1, \dots, x_T) ,
- ▶ Predict x_t for $t > T$.

(Batch) learning from data

- ▶ $x_t \in \mathcal{X}$, the **input** variables (or **features**)
- ▶ Usually $\mathcal{X} = \mathbb{R}^n$: the n -dimensional Euclidean space

Unsupervised learning:

- ▶ Training examples (x_1, \dots, x_T) ,
- ▶ Predict x_t for $t > T$.

Supervised learning:

- ▶ Given labelled training examples $(x_1, y_1), \dots, (x_T, y_T)$
- ▶ $y_t \in \mathcal{Y}$, the **output** variables (or **labels**)
- ▶ For new x_t , $t > T$, predict y_t .

(Batch) learning from data

- ▶ $x_t \in \mathcal{X}$, the **input** variables (or **features**)
- ▶ Usually $\mathcal{X} = \mathbb{R}^n$: the n -dimensional Euclidean space

Unsupervised learning:

- ▶ Training examples (x_1, \dots, x_T) ,
- ▶ Predict x_t for $t > T$.

Supervised learning:

- ▶ Given labelled training examples $(x_1, y_1), \dots, (x_T, y_T)$
- ▶ $y_t \in \mathcal{Y}$, the **output** variables (or **labels**)
- ▶ For new x_t , $t > T$, predict y_t .

Classification vs Regression

- ▶ Classification: $\mathcal{Y} = \{1, \dots, m\}$ are **discrete** labels
- ▶ Regression: $\mathcal{Y} = \mathbb{R}^m$ are **continuous** values

The kNN algorithm idea

- ▶ Assume an unknown example is similar to its neighbours
- ▶ Smoothness allows us to make predictions

Discriminatory analysis-nonparametric discrimination: consistency properties, Evelyn Fix and Joseph L. Hodges Jr, 1951.




Figure: Evelyn Fix



Figure: Joseph Hodges

Performance of KNN on image classification



- ▶ Really simple!
- ▶ Can outperform really complex models!

Introduction

The hidden secret of machine learning

The algorithm

k Nearest Neighbours

Extensions and parameters

Activities

The Nearest Neighbour algorithm

Pseudocode

- ▶ Input: Data $(x_t, y_t)_{t=1}^T$, test point x , distance d
- ▶ $t^* = \arg \min_t d(x_t, x)$ / How do we implement this?
- ▶ Return $\hat{y}_t = y_{t^*}$

Classification

$$\hat{y}_t \in [m] \equiv \{1, \dots, m\}$$

Regression

$$\hat{y}_t \in \mathbb{R}^m$$

The k -Nearest Neighbour algorithm

Pseudocode

- ▶ Input: Data $(x_t, y_t)_{t=1}^T$, test point x , distance d , neighbours k
- ▶ Calculate $h_t = d(x_t, x)$ for all t .
- ▶ Get sorted indices $s = \text{argsort}(h)$ so that $d(x_{s_i}, x) \leq d(x_{s_{i+1}}, x)$ for all i . (How?)
- ▶ Return $\sum_{i=1}^k y_{s_i} / k$.

Classification

- ▶ It is not convenient to work with discrete labels.
- ▶ We use a **one-hot encoding** $(0, \dots, 0, 1, 0, \dots, 0)$.
- ▶ $y_t \in \{0, 1\}^m$ with $\|y_t\|_1 = 1$, so that the class of the t -th example is j iff $y_{t,j} = 1$.

Regression

- ▶ $y_t \in \mathbb{R}^m$, so we need do nothing

Making a decision

kNN: A **model** of the conditional distribution $P(y|x)$

- ▶ Given features x , we get a vector
- ▶ $p_i = \hat{\mathbb{P}}(y = i|x)$.

The optimal decision rule π derived from kNN

- ▶ Classification decision $a_t \sim \pi(a|x_t)$
- ▶ $a_t \in \mathcal{A}$ but $\mathcal{A} \neq \mathcal{Y}$, e.g. can include "Do not Know", or "Alert" etc.
- ▶ Actual label y_t
- ▶ $U(a_t, y_t)$: utility function depending on the application.

Decision rule maximising accuracy

- ▶ $a_t = \arg \max_i \hat{\mathbb{P}}(y = i|x)$.

The number of neighbours

$$k = 1$$

- ▶ How does it perform on the training data?
- ▶ How might it perform on unseen data?

$$k = T$$

- ▶ How does it perform on the training data?
- ▶ How might it perform on unseen data?

Distance function

For data in \mathbb{R}^n , p -norm

$$d(x, y) = \|x - y\|_p$$

Scaled norms

When features having varying scales:

$$d(x, y) = \|Sx - Sy\|_p$$

Or pre-scale the data

Complex data

- ▶ Manifold distances
- ▶ Graph distance

Distances

A distance $d(\cdot, \cdot)$:

- ▶ Identity $d(x, x) = 0$.
- ▶ Positivity $d(x, y) > 0$ if $x \neq y$.
- ▶ Symmetry $d(y, x) = d(x, y)$.
- ▶ Triangle inequality $d(x, y) \leq d(x, z) + d(z, y)$.

For data in \mathbb{R}^n , p -norm

$$d(x, y) = \|x - y\|_p$$

Norms;

A norm $\| \cdot \|$

- ▶ Zero element $\|0\| = 0$.
- ▶ Homogeneity $\|cx\| = c\|x\|$ for any scalar a .
- ▶ Triangle inequality $\|x + y\| \leq \|x\| + \|y\|$.

ℓ_p -norm

$$\|z\|_p = \left(\sum_i z_i^p \right)^{1/p}$$

Neighbourhood calculation

If we have T datapoints

Sort and top K .

- Requires $O(T \ln T)$ time

Use the Cover-Tree or KD-Tree algorithm

- Requires $O(cK \ln T)$ time.
- c depends on the data distribution.

Introduction

The hidden secret of machine learning

The algorithm

k Nearest Neighbours

Extensions and parameters

Activities

KNN activity

- ▶ Implement nearest neighbours
- ▶ `./src/KNearestNeighbours/NearestNeighbourClassifier.py`
- ▶ Introduction to scikitlearn nearest neighbours
- ▶ Introduction to generalisation errors